

# Modern Vision Architectures

Chapter 8 · Designing depth, information flow, and reuse

Dr. A. Belcaid


## Table of contents

|          |                                                                 |           |
|----------|-----------------------------------------------------------------|-----------|
| <b>1</b> | <b>Why architecture becomes a design problem</b>                | <b>3</b>  |
| 1.1      | AlexNet as the reference point . . . . .                        | 3         |
| <b>2</b> | <b>VGG: depth through repeated small blocks</b>                 | <b>3</b>  |
| 2.1      | Why repeated $3 \times 3$ convolutions? . . . . .               | 5         |
| 2.2      | The full model becomes a short specification . . . . .          | 6         |
| <b>3</b> | <b>NiN: learn within each spatial location</b>                  | <b>6</b>  |
| 3.1      | A $1 \times 1$ convolution is a learned channel mixer . . . . . | 7         |
| 3.2      | Replace the dense head with class evidence maps . . . . .       | 7         |
| <b>4</b> | <b>Inception: branch across spatial scales</b>                  | <b>9</b>  |
| 4.1      | GoogLeNet compresses the block into a reusable symbol . . . . . | 11        |
| <b>5</b> | <b>ResNet: learn a residual correction</b>                      | <b>11</b> |
| 5.1      | Reparameterize the desired mapping . . . . .                    | 12        |
| 5.2      | Identity versus projection shortcuts . . . . .                  | 13        |
| <b>6</b> | <b>DenseNet: preserve the full feature history</b>              | <b>14</b> |
| 6.1      | The dense contract . . . . .                                    | 14        |
| 6.1.1    | Numerical growth example . . . . .                              | 15        |
| 6.2      | Transition layers cool channel growth . . . . .                 | 15        |
| 6.3      | PyTorch implementation . . . . .                                | 16        |
| <b>7</b> | <b>The architecture toolkit</b>                                 | <b>16</b> |
| <b>8</b> | <b>Compose mechanisms for a task</b>                            | <b>17</b> |
| <b>9</b> | <b>Check your understanding</b>                                 | <b>18</b> |
| 9.1      | Exercise 1: VGG receptive field and parameters . . . . .        | 18        |
| 9.2      | Exercise 2: Inception output channels . . . . .                 | 18        |
| 9.3      | Exercise 3: choose a residual shortcut . . . . .                | 18        |
| 9.4      | Exercise 4: DenseNet growth . . . . .                           | 19        |

## 10 Further reading and implementation links

19

## PART V · LEARNING FROM IMAGES



**Figure 1:** A visual architecture pipeline in which feature blocks become deeper, branch, and reconnect through a shortcut. Original course illustration created for this chapter.

Chapter 7 established the basic convolutional vocabulary: local filters, channels, nonlinearities, pooling, and receptive fields. This chapter asks the next design question: **how should those operations be organized when a vision model must become deeper, more expressive, and still trainable?**

The named architectures are historically important, but the durable objective is not to memorize five model names. It is to extract reusable mechanisms—repeat, mix channels, branch, add, and concatenate—and learn to compose them around a task specification.

**By the end of this chapter, you should be able to**

1. Explain why depth is an architectural and optimization decision, not merely a layer count.
2. Construct VGG-style repeated convolutional blocks from a compact configuration.
3. Explain how NiN uses  $1 \times 1$  convolutions and global average pooling.
4. Trace all four branches of an Inception block and account for its output channels.
5. Distinguish identity and projection shortcuts in a residual block.
6. Explain DenseNet growth rate, transition compression, and activation-memory cost.
7. Translate each architecture diagram into PyTorch code.
8. Combine mechanisms from several architecture families to satisfy one task specification.

## 1 Why architecture becomes a design problem

Consider scaling a familiar CNN from a small dataset to ImageNet-scale recognition. The input becomes larger, the number of classes grows, and the visual variation becomes much broader. Repeating the same convolution–activation–pooling recipe is tempting, but three pressures appear quickly:

| Pressure             | What can go wrong?                                | Architectural question                      |
|----------------------|---------------------------------------------------|---------------------------------------------|
| Spatial resolution   | Aggressive pooling erases small structures early. | Where should downsampling occur?            |
| Parameter allocation | Wide kernels and dense heads dominate the budget. | Where should capacity be spent?             |
| Information flow     | A deeper plain chain becomes harder to optimize.  | How can earlier signals reach later layers? |

Architecture is therefore part of the hypothesis. It decides which computations are easy to express, which signals remain available, and how gradients can travel through the model.

### 1.1 AlexNet as the reference point


AlexNet demonstrated that a deep convolutional system trained with GPUs, ReLU activations, dropout, and data augmentation could succeed on large-scale image recognition ([Krizhevsky et al. 2012](#)). Its structure also gives us a useful reference: early wide kernels and aggressive strides ingest the image, while a large fully connected head performs classification.

The architecture worked, but it left open questions. Must every layer be designed individually? Can depth be increased systematically? Can the classifier avoid spending most of the parameters? Can information follow shorter routes through a deep model?

## 2 VGG: depth through repeated small blocks

The VGG study treated architecture design as a controlled experiment ([Simonyan and Zisserman 2015](#)). Across configurations A through E, the local recipe remains stable while learned depth increases. This makes depth an experimental variable rather than an accidental consequence of unrelated layer choices.

[Read the original VGG paper](#). The authors' [Visual Geometry Group project page](#) also provides the original model information.



**Figure 2:** Course reconstruction of AlexNet as a sequence of compact layer bars. The diagram uses the visual grammar followed throughout this chapter.

Table 1: **ConvNet configurations** (shown in columns). The depth of the configurations increases from the left (A) to the right (E), as more layers are added (the added layers are shown in bold). The convolutional layer parameters are denoted as “conv(receptive field size)-(number of channels)”. The ReLU activation function is not shown for brevity.

| ConvNet Configuration       |                        |                               |                                            |                                            |                                                         |
|-----------------------------|------------------------|-------------------------------|--------------------------------------------|--------------------------------------------|---------------------------------------------------------|
| A                           | A-LRN                  | B                             | C                                          | D                                          | E                                                       |
| 11 weight layers            | 11 weight layers       | 13 weight layers              | 16 weight layers                           | 16 weight layers                           | 19 weight layers                                        |
| input (224 × 224 RGB image) |                        |                               |                                            |                                            |                                                         |
| conv3-64                    | conv3-64<br><b>LRN</b> | conv3-64<br><b>conv3-64</b>   | conv3-64<br>conv3-64                       | conv3-64<br>conv3-64                       | conv3-64<br>conv3-64                                    |
| maxpool                     |                        |                               |                                            |                                            |                                                         |
| conv3-128                   | conv3-128              | conv3-128<br><b>conv3-128</b> | conv3-128<br>conv3-128                     | conv3-128<br>conv3-128                     | conv3-128<br>conv3-128                                  |
| maxpool                     |                        |                               |                                            |                                            |                                                         |
| conv3-256<br>conv3-256      | conv3-256<br>conv3-256 | conv3-256<br>conv3-256        | conv3-256<br>conv3-256<br><b>conv1-256</b> | conv3-256<br>conv3-256<br><b>conv3-256</b> | conv3-256<br>conv3-256<br>conv3-256<br><b>conv3-256</b> |
| maxpool                     |                        |                               |                                            |                                            |                                                         |
| conv3-512<br>conv3-512      | conv3-512<br>conv3-512 | conv3-512<br>conv3-512        | conv3-512<br>conv3-512<br><b>conv1-512</b> | conv3-512<br>conv3-512<br><b>conv3-512</b> | conv3-512<br>conv3-512<br>conv3-512<br><b>conv3-512</b> |
| maxpool                     |                        |                               |                                            |                                            |                                                         |
| conv3-512<br>conv3-512      | conv3-512<br>conv3-512 | conv3-512<br>conv3-512        | conv3-512<br>conv3-512<br><b>conv1-512</b> | conv3-512<br>conv3-512<br><b>conv3-512</b> | conv3-512<br>conv3-512<br>conv3-512<br><b>conv3-512</b> |
| maxpool                     |                        |                               |                                            |                                            |                                                         |
| FC-4096                     |                        |                               |                                            |                                            |                                                         |
| FC-4096                     |                        |                               |                                            |                                            |                                                         |
| FC-1000                     |                        |                               |                                            |                                            |                                                         |
| soft-max                    |                        |                               |                                            |                                            |                                                         |


**Figure 3:** Original VGG configuration table showing networks A through E. Source: Simonyan and Zisserman, Table 1.

## 2.1 Why repeated $3 \times 3$ convolutions?

Two padded  $3 \times 3$  convolutions have an effective receptive field of  $5 \times 5$ ; three have an effective receptive field of  $7 \times 7$ . Stacking small kernels has two advantages:

1. it inserts a nonlinearity after every convolution;
2. it often uses fewer parameters than one equally wide large kernel.

For equal input and output width  $C$ , one  $5 \times 5$  convolution uses approximately  $25C^2$  weights. Two  $3 \times 3$  convolutions use approximately  $18C^2$  weights. The comparison is not exact when intermediate channel widths differ, but it explains the design pressure.



**Figure 4:** A VGG block groups repeated padded three-by-three convolutions and ReLU operations before one max-pooling step. Original course SVG.

The block preserves  $H \times W$  while extracting features, then pooling halves the spatial resolution. This ordering matters: the model performs several learned transformations before discarding spatial detail.

```
import torch
from torch import nn

def vgg_block(num_convs: int, out_channels: int) -> nn.Sequential:
    layers = []
```

```

for _ in range(num_convs):
    layers.extend([
        nn.LazyConv2d(out_channels, kernel_size=3, padding=1),
        nn.ReLU(),
    ])
layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
return nn.Sequential(*layers)

```

The [Conv2d documentation](#) defines the channel, kernel, stride, and padding contract. `Sequential` is appropriate because each operation has exactly one successor.

## 2.2 The full model becomes a short specification

```

class VGG(nn.Module):
    def __init__(self, arch, num_classes=10):
        super().__init__()
        blocks = [vgg_block(n, c) for n, c in arch]
        self.net = nn.Sequential(
            *blocks,
            nn.Flatten(),
            nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
            nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
            nn.LazyLinear(num_classes),
        )

    def forward(self, x):
        return self.net(x)


vgg11 = ((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))
model = VGG(vgg11)

```

VGG's contribution is methodological: standardize the local rule, then express the architecture as a sequence of block settings. Its principal drawbacks are the expensive dense classifier, high convolutional cost, and a plain information path through increasing depth.

## 3 NiN: learn within each spatial location

Network in Network (NiN) responds to two VGG-era costs: limited channel interaction inside a convolutional stage and a very large dense classifier ([Lin et al. 2014](#)). [Read the original NiN paper](#).



**Figure 5:** VGG represented as a full configuration, a reusable VGG block, and the AlexNet reference architecture. Original course SVG.

### 3.1 A $1 \times 1$ convolution is a learned channel mixer

At a fixed spatial position  $(h, w)$ , a  $1 \times 1$  convolution computes

$$y_{o,h,w} = b_o + \sum_{c=1}^{C_{in}} W_{o,c} x_{c,h,w}.$$


It does not collect neighboring pixels. It learns a new combination of the channels already available at that location. With an activation after each  $1 \times 1$  convolution, the block becomes a small multilayer perceptron shared across all spatial positions.

```
def nin_block(out_channels, kernel_size, stride, padding):
    return nn.Sequential(
        nn.LazyConv2d(out_channels, kernel_size, stride, padding),
        ↪ nn.ReLU(),
        nn.LazyConv2d(out_channels, kernel_size=1), nn.ReLU(),
        nn.LazyConv2d(out_channels, kernel_size=1), nn.ReLU(),
    )
```

### 3.2 Replace the dense head with class evidence maps


NiN's final block emits `num_classes` channels. Channel  $k$  becomes a spatial evidence map for class  $k$ . Global average pooling then converts each map into one scalar:

$$s_k = \frac{1}{HW} \sum_{h=1}^H \sum_{w=1}^W A_{k,h,w}.$$



**Figure 6:** A NiN block begins with a spatial convolution and follows it with two one-by-one convolutional channel mixers. Original course SVG.

This gives the head an interpretable contract: one map per class, one average per map. PyTorch's `AdaptiveAvgPool2d` can force the output to  $1 \times 1$  regardless of the incoming spatial size.



**Figure 7:** Full NiN architecture with four grouped NiN blocks and a global-average-pooling classifier. Original course SVG.

```
class NiN(nn.Module):
    def __init__(self, num_classes=10):
        super().__init__()
        self.net = nn.Sequential(
            nin_block(96, 11, 4, 0), nn.MaxPool2d(3, stride=2),
            nin_block(256, 5, 1, 2), nn.MaxPool2d(3, stride=2),
            nin_block(384, 3, 1, 1), nn.MaxPool2d(3, stride=2),
            nn.Dropout(0.5),
            nin_block(num_classes, 3, 1, 1),
            nn.AdaptiveAvgPool2d((1, 1)),
            nn.Flatten(),
        )


    def forward(self, x):
        return self.net(x)
```

NiN makes channel mixing explicit and removes the dense head, but its early stages still commit to one kernel size and one sequential route.

## 4 Inception: branch across spatial scales

An Inception block asks several spatial questions in parallel. Instead of choosing one receptive-field size in advance, it evaluates  $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$ , and pooled evidence,

then concatenates their outputs along the channel dimension (Szegedy et al. 2015). [Read the original GoogLeNet paper.](#)



**Figure 8:** An Inception block with four parallel branches whose outputs are concatenated along channels. Original course SVG.

If every branch preserves  $H \times W$ , concatenation is legal and the output channel count is

$$C_{out} = c_1 + c_{22} + c_{32} + c_4.$$

The  $1 \times 1$  operations before the wider kernels are capacity valves. A direct  $5 \times 5$  convolution from  $C_{in}$  to  $C_{out}$  uses  $25C_{in}C_{out}$  weights. Reducing first to  $C_r$  channels changes the approximate cost to

$$C_{in}C_r + 25C_rC_{out}.$$

When  $C_r \ll C_{in}$ , the saving can be substantial.


```
class Inception(nn.Module):
    def __init__(self, c1, c2, c3, c4):
        super().__init__()
        self.b1 = nn.Sequential(
            nn.LazyConv2d(c1, 1), nn.ReLU())
        self.b2 = nn.Sequential(
            nn.LazyConv2d(c2[0], 1), nn.ReLU(),
            nn.LazyConv2d(c2[1], 3, padding=1), nn.ReLU())
        self.b3 = nn.Sequential(
            nn.LazyConv2d(c3[0], 1), nn.ReLU(),
            nn.LazyConv2d(c3[1], 5, padding=2), nn.ReLU())
        self.b4 = nn.Sequential(
            nn.MaxPool2d(3, stride=1, padding=1),
            nn.LazyConv2d(c4, 1), nn.ReLU())
```

```
def forward(self, x):
    return torch.cat(
        [self.b1(x), self.b2(x), self.b3(x), self.b4(x)], dim=1)
```

`torch.cat` concatenates tensors along an existing dimension. For NCHW images, `dim=1` is the channel dimension. Batch size, height, and width must agree across branches.

#### 4.1 GoogLeNet compresses the block into a reusable symbol

Once the internal four-branch contract is understood, a full network can treat each Inception block as one module. GoogLeNet uses nine such blocks in groups of two, five, and two, with pooling between groups and global average pooling in the head.




**Figure 9:** Compressed GoogLeNet architecture assembled from nine Inception modules. Original course SVG.

The architecture gains multi-scale processing and parameter-efficient wide branches. The cost is specification complexity: every branch requires channel choices, parallel execution can be hardware-sensitive, and the original model used auxiliary classifiers during training.

## 5 ResNet: learn a residual correction

Adding layers to a plain network does not guarantee lower training error. The original ResNet paper compared 20- and 56-layer plain networks on CIFAR-10; the deeper plain network had *higher training error* (He et al. 2016). This is an optimization failure, not merely overfitting.

[Read the original ResNet paper](#) or the [CVPR open-access version](#).



**Figure 10:** Training and test error curves for plain 20- and 56-layer networks. The deeper plain network performs worse. Source: He et al., Figure 1.

### 5.1 Reparameterize the desired mapping


Suppose the desired block mapping is  $H(x)$ . A plain stack attempts to learn  $H$  directly. A residual block instead learns

$$F(x) = H(x) - x,$$

and returns

$$y = F(x) + x.$$

If the best transformation is close to identity, setting  $F(x) \approx 0$  is easier than forcing several nonlinear layers to reproduce  $x$  exactly.




**Figure 11:** Residual block with a learned two-convolution path and an identity shortcut meeting at addition. Original course SVG.

The shortcut also creates a direct route for information and gradients. It does not eliminate every optimization difficulty, but it changes the geometry of the problem and makes identity preservation structurally available.

## 5.2 Identity versus projection shortcuts

Elementwise addition requires equal shapes. When the residual branch preserves  $(H, W, C)$ , the shortcut should normally remain the free identity. At a stage boundary, the main path may halve spatial resolution and increase channels. A  $1 \times 1$  projection applies the same stride and maps the shortcut to the required channel count.



**Figure 12:** Identity and projection residual shortcuts compared side by side. Original course SVG.

```
class ResidualBlock(nn.Module):
    def __init__(self, in_ch, out_ch, stride=1, use_projection=False):
        super().__init__()
        self.main = nn.Sequential(
            nn.Conv2d(in_ch, out_ch, 3, stride,
                      padding=1, bias=False),
            nn.BatchNorm2d(out_ch), nn.ReLU(),
            nn.Conv2d(out_ch, out_ch, 3,
                      padding=1, bias=False),
            nn.BatchNorm2d(out_ch),
        )
        self.shortcut = (
            nn.Conv2d(in_ch, out_ch, 1, stride, bias=False)
            if use_projection else None
        )

    def forward(self, x):
        identity = x if self.shortcut is None else self.shortcut(x)
        return torch.relu(self.main(x) + identity)
```


**BatchNorm2d** normalizes each channel using batch statistics during training and running estimates during evaluation. Its placement differs in later pre-activation ResNets; this chapter follows the original basic post-activation teaching block.

### 💡 Shape audit

Before every residual addition, write the shapes of both operands. If either spatial resolution or channel count differs, an identity shortcut cannot be added directly.

## 6 DenseNet: preserve the full feature history

ResNet combines a learned correction and a preserved signal by addition. DenseNet makes a different choice: each layer receives the concatenation of *all* preceding feature maps, and its output remains available to all subsequent layers (Huang et al. 2017).



**Figure 1:** A 5-layer dense block with a growth rate of  $k = 4$ . Each layer takes all preceding feature-maps as input.

**Figure 13:** Original DenseNet paper Figure 1 showing a five-layer dense block with growth rate four.

Read the [original DenseNet paper](#), the [CVPR open-access paper](#), or the authors' [reference implementation](#).

### 6.1 The dense contract

For layer  $\ell$ , define the collective input

$$z_\ell = [x_0, x_1, \dots, x_{\ell-1}],$$


where brackets denote channel concatenation. The layer contributes

$$x_\ell = H_\ell(z_\ell).$$

If the block begins with  $C_0$  channels and every layer contributes  $k$  new maps, then after  $L$  layers

$$C_L = C_0 + Lk.$$

The value  $k$  is the **growth rate**. Dense layers can remain narrow because they do not need to recreate useful earlier features.



**Figure 14:** Four dense layers append new features to a shared feature history. Original course SVG.

### 6.1.1 Numerical growth example

Let  $C_0 = 8$ ,  $k = 4$ , and  $L = 4$ . The channel counts after successive layers are

$$8 \rightarrow 12 \rightarrow 16 \rightarrow 20 \rightarrow 24.$$

The block contains 24 output channels, but the fourth layer itself computes only four new maps.

## 6.2 Transition layers cool channel growth

Dense concatenation widens the representation continuously. A transition layer commonly applies a  $1 \times 1$  convolution followed by  $2 \times 2$  average pooling. With compression factor  $0 < \theta \leq 1$ , the transition emits

$$C_{transition} = \lfloor \theta C \rfloor$$

channels and halves spatial resolution. Compression controls capacity and memory; it should not be confused with regularization. The paper also uses dropout in some small-dataset experiments to reduce feature co-adaptation.

### 6.3 PyTorch implementation

```
class DenseLayer(nn.Module):
    def __init__(self, in_ch, growth_rate):
        super().__init__()
        self.transform = nn.Sequential(
            nn.BatchNorm2d(in_ch), nn.ReLU(),
            nn.Conv2d(in_ch, growth_rate, 3,
                      padding=1, bias=False),
        )

    def forward(self, x):
        new_features = self.transform(x)
        return torch.cat([x, new_features], dim=1)

def dense_block(in_ch, num_layers, growth_rate):
    layers = []
    for i in range(num_layers):
        current_channels = in_ch + i * growth_rate
        layers.append(DenseLayer(current_channels, growth_rate))
    return nn.Sequential(*layers)
```

For `dense_block(64, 4, 32)`, the output width is

$$64 + 4(32) = 192 \text{ channels.}$$

DenseNet creates short routes, encourages feature reuse, and can be parameter-efficient. Its principal systems cost is activation memory and data movement: earlier maps must remain available, and repeated concatenation is not free.

## 7 The architecture toolkit

The chapter's five families can be summarized as operations rather than model names.

| Family    | Reusable operation                                 | Pressure addressed                                | Main tradeoff                        |
|-----------|----------------------------------------------------|---------------------------------------------------|--------------------------------------|
| VGG       | <b>Repeat</b> a uniform small-kernel block         | Systematic depth                                  | High convolutional cost              |
| NiN       | <b>Mix channels</b> with $1 \times 1$ convolutions | Local cross-channel interaction and compact heads | Rigid class-map head                 |
| Inception | <b>Branch</b> across receptive-field sizes         | Multi-scale evidence                              | Hand-tuned branch complexity         |
| ResNet    | <b>Add</b> a correction to an identity path        | Stable deep optimization                          | Addition requires shape alignment    |
| DenseNet  | <b>Concatenate</b> the full feature history        | Explicit feature reuse and short paths            | Activation memory and channel growth |

These mechanisms are compatible. A modern model may use repeated stages, bottleneck  $1 \times 1$  convolutions, parallel branches, residual additions, and selective concatenation in the same architecture.

## 8 Compose mechanisms for a task

Suppose the design brief is:

Classify ten plant diseases from mobile photographs. Lesions may be small, compute is limited, and the network must be deep enough to learn complex texture while remaining trainable.

A defensible first design could combine:

1. **VGG's repetition:** a regular stem of padded  $3 \times 3$  convolutions.
2. **Inception's branching:** a multi-scale stage for lesions of different apparent sizes.
3. **ResNet's addition:** identity shortcuts through deeper stages.
4. **NiN's head:** class maps followed by global average pooling.

Dense concatenation is not intrinsically wrong for this task. It is omitted from this first proposal because its activation-memory cost conflicts with the mobile constraint. A different deployment budget could change that decision.

The design loop is:

specification  $\rightarrow$  pressure  $\rightarrow$  mechanism  $\rightarrow$  evidence.

Architecture names help us trace ideas historically. A model design, however, should be justified as a composition of mechanisms chosen for a task.

## 9 Check your understanding

### 9.1 Exercise 1: VGG receptive field and parameters

Two consecutive convolutions use  $3 \times 3$  kernels, stride one, padding one, and 64 input/output channels. Determine their effective receptive field and compare their weight count with one  $5 \times 5$  convolution from 64 to 64 channels. Ignore biases.

Worked solution

Two stride-one  $3 \times 3$  convolutions produce an effective  $5 \times 5$  receptive field. Their weight count is

$$2(3 \cdot 3 \cdot 64 \cdot 64) = 73,728.$$

One  $5 \times 5$  convolution uses

$$5 \cdot 5 \cdot 64 \cdot 64 = 102,400$$

weights. The stacked version uses fewer weights and inserts an additional nonlinearity.

### 9.2 Exercise 2: Inception output channels

An Inception block emits 64 channels from its direct  $1 \times 1$  branch, 128 from its  $3 \times 3$  branch, 32 from its  $5 \times 5$  branch, and 32 from its pooling branch. What is the output shape for input  $(N, 192, 28, 28)$  if every branch preserves resolution?

Worked solution

Concatenation adds channel counts:

$$C_{out} = 64 + 128 + 32 + 32 = 256.$$

The output shape is  $(N, 256, 28, 28)$ .

### 9.3 Exercise 3: choose a residual shortcut

For each transition, decide whether the shortcut can be identity or requires projection.

1.  $(N, 64, 32, 32)$  to  $(N, 64, 32, 32)$
2.  $(N, 64, 32, 32)$  to  $(N, 128, 16, 16)$

Worked solution

1. The shapes agree, so an identity shortcut is valid.

2. Channels and spatial dimensions change. Use a  $1 \times 1$  projection with stride two to produce  $(N, 128, 16, 16)$  before addition.

#### 9.4 Exercise 4: DenseNet growth

A dense block begins with 96 channels, contains six layers, and uses growth rate  $k = 24$ . Find its output channels. If a transition uses  $\theta = 0.5$ , how many channels remain after compression?

Worked solution

The dense block emits

$$C_6 = 96 + 6(24) = 240$$

channels. Compression gives

$$\lfloor 0.5(240) \rfloor = 120$$

channels.

### 10 Further reading and implementation links

- [AlexNet: ImageNet Classification with Deep Convolutional Neural Networks](#)
- [VGG: Very Deep Convolutional Networks for Large-Scale Image Recognition](#)
- [Network in Network](#)
- [GoogLeNet: Going Deeper with Convolutions](#)
- [ResNet: Deep Residual Learning for Image Recognition](#)
- [DenseNet: Densely Connected Convolutional Networks](#)
- [PyTorch nn.Module](#)
- [PyTorch nn.Sequential](#)
- [PyTorch nn.ModuleList](#)
- [PyTorch Conv2d](#)
- [PyTorch BatchNorm2d](#)
- [PyTorch MaxPool2d](#)
- [PyTorch AdaptiveAvgPool2d](#)
- [PyTorch Dropout](#)
- [PyTorch torch.cat](#)
- [PyTorch tutorial: What is torch.nn really?](#)

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. “Deep Residual Learning for Image Recognition.” *IEEE Conference on Computer Vision and Pattern Recognition*, 770–78. <https://doi.org/10.1109/CVPR.2016.90>.

- Huang, Gao, Zhuang Liu, Laurens van der Maaten, and Kilian Q. Weinberger. 2017. “Densely Connected Convolutional Networks.” *IEEE Conference on Computer Vision and Pattern Recognition*, 4700–4708. <https://doi.org/10.1109/CVPR.2017.243>.
- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. 2012. “ImageNet Classification with Deep Convolutional Neural Networks.” *Advances in Neural Information Processing Systems* 25. <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>.
- Lin, Min, Qiang Chen, and Shuicheng Yan. 2014. “Network in Network.” *International Conference on Learning Representations*. <https://arxiv.org/abs/1312.4400>.
- Simonyan, Karen, and Andrew Zisserman. 2015. “Very Deep Convolutional Networks for Large-Scale Image Recognition.” *International Conference on Learning Representations*. <https://arxiv.org/abs/1409.1556>.
- Szegedy, Christian, Wei Liu, Yangqing Jia, et al. 2015. “Going Deeper with Convolutions.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 1–9. <https://arxiv.org/abs/1409.4842>.